

Backpropagation: The Good, the Bad and the Ugly

Keith L. Downing


The Norwegian University of Science and Technology (NTNU)
Trondheim, Norway
keithd@idi.ntnu.no

January 16, 2023

Supervised Learning


- Constant feedback from an instructor, indicating not only right/wrong, but also the **correct** answer for each training case.
- Many cases (i.e., input-output pairs) to be learned.
- Weights are modified by a complex procedure (back-propagation) based on output error.
- Feed-forward networks with back-propagation learning are the standard implementation.
- 99% of neural network applications use this.
- Typical usage: problems with a) lots of input-output training data, and b) goal of a mapping (function) from inputs to outputs.
- **Not** biologically plausible, although the cerebellum appears to exhibit some aspects.
- But, the **result** of backprop, a trained ANN to perform some function, can be very useful to neuroscientists as a **sufficiency proof**.

Backpropagation Overview




- **Feed-Forward Phase** - Inputs sent through the ANN to compute outputs.
- **Feedback Phase** - Error passed back from output to input layers and used to update weights along the way.

Training -vs- Testing



- **Generalization** - correctly handling test cases (that ANN has not been trained on).
- **Over-Training** - weights become so fine-tuned to the training cases that generalization suffers: failure on many test cases.


Widrow-Hoff (a.k.a. Delta) Rule




- Delta (δ) = error; Eta (η) = learning rate
- Goal: change w so as to reduce $|\delta|$.
- Intuitive: If $\delta > 0$, then we want to decrease it, so we must increase Y . Thus, we must increase the sum of weighted inputs to N , and we do that by increasing (decreasing) w if X is positive (negative).
- Similar for $\delta < 0$
- Assumes derivative of N 's transfer function is everywhere non-negative.

Gradient Descent

- Goal = minimize total error across all output nodes
- Method = modify weights throughout the network (i.e., at all levels) to follow the route of steepest descent in error space.



$$\Delta w_{ij} = -\eta \frac{\partial E_i}{\partial w_{ij}}$$




Sum of Squared Errors (SSE)

$$E_i = \frac{1}{2} \sum_{d \in D} (t_{id} - o_{id})^2$$

$$\frac{\partial E}{\partial w_{ij}} = \frac{1}{2} \sum_{d \in D} 2(t_{id} - o_{id}) \frac{\partial (t_{id} - o_{id})}{\partial w_{ij}} = \sum_{d \in D} (t_{id} - o_{id}) \frac{\partial (-o_{id})}{\partial w_{ij}}$$

Computing $\frac{\partial(-o_{id})}{\partial w_{ij}}$



Since output = $f(\text{sum weighted inputs})$

$$\frac{\partial E}{\partial w_{ij}} = \sum_{d \in D} (t_{id} - o_{id}) \frac{\partial(-f_T(\text{sum}_{id}))}{\partial w_{ij}}$$

where

$$\text{sum}_{id} = \sum_{k=1}^n w_{ik} x_{kd}$$

Using Chain Rule: $\frac{\partial f(g(x))}{\partial x} = \frac{\partial f}{\partial g(x)} \times \frac{\partial g(x)}{\partial x}$

$$\frac{\partial(f_T(\text{sum}_{id}))}{\partial w_{ij}} = \frac{\partial f_T(\text{sum}_{id})}{\partial \text{sum}_{id}} \times \frac{\partial \text{sum}_{id}}{\partial w_{ij}} = \frac{\partial f_T(\text{sum}_{id})}{\partial \text{sum}_{id}} \times x_{jd}$$

Computing $\frac{\partial sum_{id}}{\partial w_{ij}}$ - Easy!!

$$\begin{aligned}\frac{\partial sum_{id}}{\partial w_{ij}} &= \frac{\partial (\sum_{k=1}^n w_{ik} x_{kd})}{\partial w_{ij}} = \frac{\partial (w_{i1} x_{1d} + w_{i2} x_{2d} + \dots + w_{ij} x_{jd} + \dots + w_{in} x_{nd})}{\partial w_{ij}} \\ &= \frac{\partial (w_{i1} x_{1d})}{\partial w_{ij}} + \frac{\partial (w_{i2} x_{2d})}{\partial w_{ij}} + \dots + \frac{\partial (w_{ij} x_{jd})}{\partial w_{ij}} + \dots + \frac{\partial (w_{in} x_{nd})}{\partial w_{ij}} \\ &= 0 + 0 + \dots + x_{jd} + \dots + 0 = x_{jd}\end{aligned}$$

Computing $\frac{\partial f_T(\text{sum}_{id})}{\partial \text{sum}_{id}}$ - Harder for some f_T

f_T = Identity function: $f_T(\text{sum}_{id}) = \text{sum}_{id}$

$$\frac{\partial f_T(\text{sum}_{id})}{\partial \text{sum}_{id}} = 1$$

Thus:

$$\frac{\partial (f_T(\text{sum}_{id}))}{\partial w_{ij}} = \frac{\partial f_T(\text{sum}_{id})}{\partial \text{sum}_{id}} \times \frac{\partial \text{sum}_{id}}{\partial w_{ij}} = 1 \times x_{jd} = x_{jd}$$

f_T = Sigmoid: $f_T(\text{sum}_{id}) = \frac{1}{1 + e^{-\text{sum}_{id}}}$

$$\frac{\partial f_T(\text{sum}_{id})}{\partial \text{sum}_{id}} = o_{id}(1 - o_{id})$$

Thus:

$$\frac{\partial (f_T(\text{sum}_{id}))}{\partial w_{ij}} = \frac{\partial f_T(\text{sum}_{id})}{\partial \text{sum}_{id}} \times \frac{\partial \text{sum}_{id}}{\partial w_{ij}} = o_{id}(1 - o_{id}) \times x_{jd} = o_{id}(1 - o_{id})x_{jd}$$

The only non-trivial calculation

$$\begin{aligned}\frac{\partial f_T(\text{sum}_{id})}{\partial \text{sum}_{id}} &= \frac{\partial ((1 + e^{-\text{sum}_{id}})^{-1})}{\partial \text{sum}_{id}} = (-1) \frac{\partial (1 + e^{-\text{sum}_{id}})}{\partial \text{sum}_{id}} (1 + e^{-\text{sum}_{id}})^{-2} \\ &= (-1)(-1)e^{-\text{sum}_{id}} (1 + e^{-\text{sum}_{id}})^{-2} = \frac{e^{-\text{sum}_{id}}}{(1 + e^{-\text{sum}_{id}})^2}\end{aligned}$$

But notice that:

$$\frac{e^{-\text{sum}_{id}}}{(1 + e^{-\text{sum}_{id}})^2} = f_T(\text{sum}_{id})(1 - f_T(\text{sum}_{id})) = o_{id}(1 - o_{id})$$

Putting it all together

$$\frac{\partial E_i}{\partial w_{ij}} = \sum_{d \in D} (t_{id} - o_{id}) \frac{\partial (-f_T(\text{sum}_{id}))}{\partial w_{ij}} = - \sum_{d \in D} \left((t_{id} - o_{id}) \frac{\partial f_T(\text{sum}_{id})}{\partial \text{sum}_{id}} \times \frac{\partial \text{sum}_{id}}{\partial w_{ij}} \right)$$

So for $f_T = \text{Identity}$:

$$\frac{\partial E_i}{\partial w_{ij}} = - \sum_{d \in D} (t_{id} - o_{id}) x_{jd}$$

and for $f_T = \text{Sigmoid}$:

$$\frac{\partial E_i}{\partial w_{ij}} = - \sum_{d \in D} (t_{id} - o_{id}) o_{id} (1 - o_{id}) x_{jd}$$

Weight Updates ($f_T = \text{Sigmoid}$)

Batch: update weights after each training **epoch**

$$\Delta w_{ij} = -\eta \frac{\partial E_i}{\partial w_{ij}} = \eta \sum_{d \in D} (t_{id} - o_{id}) o_{id} (1 - o_{id}) x_{jd}$$

The weight changes are actually computed after each training case, but w_{ij} is not updated until the epoch's end.

Incremental: update weights after each training **case**

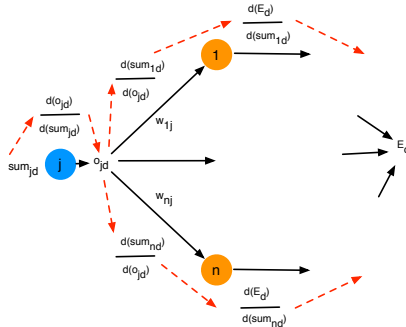
$$\Delta w_{ij} = -\eta \frac{\partial E_i}{\partial w_{ij}} = \eta (t_{id} - o_{id}) o_{id} (1 - o_{id}) x_{jd}$$

- A lower learning rate (η) recommended here than for Batch method.
- Can be dependent upon case-presentation order. So randomly sort the cases after each epoch.

Stochastic Gradient Descent (SGD)

Similar to Batch, but processing (feed forward + gradient calculations) a subset (minibatch) of D between each weight update.

Backpropagation in Multi-Layered Neural Networks




For each node (j) and each training case (d), backpropagation computes an error term:

$$\delta_{jd} = \frac{\partial E_d}{\partial sum_{jd}}$$

by calculating the influence of sum_{jd} along each connection from node j to the next downstream layer.

Computing δ_{jd}



Along the upper path, the contribution to $\frac{\partial E_d}{\partial sum_{jd}}$ is:

$$\frac{\partial o_{jd}}{\partial sum_{jd}} \times \frac{\partial sum_{1d}}{\partial o_{jd}} \times \frac{\partial E_d}{\partial sum_{1d}}$$

So summing along all paths:

$$\frac{\partial E_d}{\partial sum_{jd}} = \frac{\partial o_{jd}}{\partial sum_{jd}} \sum_{k=1}^n \frac{\partial sum_{kd}}{\partial o_{jd}} \frac{\partial E_d}{\partial sum_{kd}}$$

Computing δ_{jd}

Just as before, most terms are 0 in the derivative of the sum, so:

$$\frac{\partial \text{sum}_{kd}}{\partial o_{jd}} = w_{kj}$$

Assuming f_T = a sigmoid:

$$\frac{\partial o_{jd}}{\partial \text{sum}_{jd}} = \frac{\partial f_T(\text{sum}_{jd})}{\partial \text{sum}_{jd}} = o_{jd}(1 - o_{jd})$$

Thus:


$$\begin{aligned}\delta_{jd} &= \frac{\partial E_d}{\partial \text{sum}_{jd}} = \frac{\partial o_{jd}}{\partial \text{sum}_{jd}} \sum_{k=1}^n \frac{\partial \text{sum}_{kd}}{\partial o_{jd}} \frac{\partial E_d}{\partial \text{sum}_{kd}} \\ &= o_{jd}(1 - o_{jd}) \sum_{k=1}^n w_{kj} \delta_{kd}\end{aligned}$$

Note that δ_{jd} is defined recursively in terms of the δ values in the next downstream layer:

$$\delta_{jd} = o_{jd}(1 - o_{jd}) \sum_{k=1}^n w_{kj} \delta_{kd}$$

So all δ values in the network can be computed by moving backwards, one layer at a time.

Computing $\frac{\partial E_d}{\partial w_{ij}}$ from δ_{jd} - Easy!!



The only effect of w_{ij} upon the error is via its effect upon sum_{id} , which is:

$$\frac{\partial sum_{id}}{\partial w_{ij}} = o_{jd}$$

So:

$$\frac{\partial E_d}{\partial w_{ij}} = \frac{\partial sum_{id}}{\partial w_{ij}} \times \frac{\partial E_d}{\partial sum_{id}} = \frac{\partial sum_{id}}{\partial w_{ij}} \times (\delta_{id}) = o_{jd} \delta_{id}$$




Computing Δw_{ij}

Given an error term, δ_{id} (for node i on training case d), the update of w_{ij} for all nodes j that feed into i is:

$$\Delta w_{ij} = -\eta \frac{\partial E_d}{\partial w_{ij}} = -\eta(o_{jd}\delta_{id}) = -\eta\delta_{id}o_{jd}$$

So given δ_{id} , you can easily calculate Δw_{ij} for all incoming arcs.


Learning XOR




- Epoch = All 4 entries of the XOR truth table.
- 2 (inputs) X 2 (hidden) X 1 (output) network
- Random init of all weights in $[-1, 1]$.
- Not linearly separable, so it takes awhile!
- Each run is different due to random weight init.

Learning to Classify Wines


Class	Properties							
1	14.23	1.71	2.43	15.6	127	2.8
1	13.2	1.78	2.14	11.2	100	2.65
2	13.11	1.01	1.7	15	78	2.98
3	13.17	2.59	2.37	20	120	1.65
⋮								




Wine Runs




13x5x1 (lr = 0.3)



13x5x1 (lr = 0.1)



13 x 10 x 1 (lr = 0.3)



13 x 25 x 1 (lr = 0.3)


*I did everything he did, but
backwards and in high heels*

.... Ginger Rogers describing her career as Fred Astaire's dance partner.

Now we'll go through similar backpropagation derivations, but with matrices: backwards in high heels.

Matrix Notation

- In the literature, the notation w_{ij} may EITHER denote the weight on the connection from node j to node i OR the connection from node i to node j . This varies; there is no standard.
- When working with matrices, it is often easier when $w_{ij} :=$ weight on connection $i \rightarrow j$ **This is the case in the remainder of this slide series.**
- Then, the normal syntax for matrix row-column indices implies that:
 - The values in row i are written as $w_{i?}$ and denote the weights on the **outputs** of node i .
 - The values in column j are written as $w_{?j}$ and denote the weights on the **inputs** to node j .



Tensor Operations (Notation, Rules)

- Tensors = arrays of 1 or more dimensions. They will be represented by CAPITAL letters. Scalars will be in small letters, often with subscripts.
- $W \bullet X$ = dot product of vectors OR the multiplication of a matrix with another matrix or vector, depending upon the situation. This is similar to the semantics of the (very versatile) **numpy.dot** function.
- The **Product Rule** for tensors (A, B, X):

$$\frac{\partial(A \bullet B)}{\partial X} = A \bullet \frac{\partial B}{\partial X} + \frac{\partial A}{\partial X} \bullet B$$

- The **Chain Rule** for tensor X and functions F and G, which take tensors as input and produce tensors as output.

$$\frac{\partial F(G(X))}{\partial X} = \frac{\partial F(.)}{\partial G(.)} \bullet \frac{\partial G(X)}{\partial X}$$

- If tensors A and B have dimensions m and n, respectively, then the Jacobian $J = \frac{\partial A}{\partial B}$ has dimensions m x n.

Tensor Tension

- In the **numerator layout** for Jacobians:

$$J_{i,\dots q,r,\dots z} = \frac{\partial A_{i,\dots q}}{\partial B_{r,\dots z}}$$

- In the **denominator layout** for Jacobians:

$$J_{i,\dots q,r,\dots z} = \frac{\partial A_{r,\dots z}}{\partial B_{i,\dots q}}$$

- The numerator layout is more common but is not a hard standard.
- These slides will use the numerator layout **as much as possible**. Any usage of denominator layout will be noted.
- Tensor calculus is the combination of linear algebra and multivariate calculus. It is very useful for mathematically describing the operations of deep-learning systems, but occasionally the link between theory and implementation gets **twisted** a bit.
- I try to minimize these twists, and I mark these twists/hacks.
- For more on tensor calculus, go to: explained.ai/matrix-calculus/

Numerator -vs- Denominator Layout


- **Numerator Layout** (A's index remains constant across each row; the first dimension of J pertains to A.)

$$J_B^A = \begin{bmatrix} \frac{\partial a_1}{\partial b_1} & \frac{\partial a_1}{\partial b_2} & \frac{\partial a_1}{\partial b_3} \\ \frac{\partial a_2}{\partial b_1} & \frac{\partial a_2}{\partial b_2} & \frac{\partial a_2}{\partial b_3} \\ \frac{\partial a_3}{\partial b_1} & \frac{\partial a_3}{\partial b_2} & \frac{\partial a_3}{\partial b_3} \end{bmatrix} \quad (1)$$

- **Denominator Layout** (B's index remains constant across each row; the first dimension of J pertains to B.)

$$J_B^A = \begin{bmatrix} \frac{\partial a_1}{\partial b_1} & \frac{\partial a_2}{\partial b_1} & \frac{\partial a_3}{\partial b_1} \\ \frac{\partial a_1}{\partial b_2} & \frac{\partial a_2}{\partial b_2} & \frac{\partial a_3}{\partial b_2} \\ \frac{\partial a_1}{\partial b_3} & \frac{\partial a_2}{\partial b_3} & \frac{\partial a_3}{\partial b_3} \end{bmatrix} \quad (2)$$

Backpropagation using Tensors



$$R = XV$$

$$Y = g(R)$$

$$S = YW$$

$$Z = f(S)$$

$$X = [x_1, x_2, x_3]$$

$$Y = [y_1, y_2]$$

$$Z = [z_1]$$

$$V = \begin{bmatrix} v_{11} & v_{12} \\ v_{21} & v_{22} \\ v_{31} & v_{32} \end{bmatrix}$$

$$W = \begin{bmatrix} w_{11} \\ w_{21} \end{bmatrix}$$

$$f = g = \text{sigmoid}$$

v_{ij} = weight from i to j

Goal: Compute $\frac{\partial Z}{\partial W}$

- Using the Chain Rule of Tensor Calculus:

$$\frac{\partial Z}{\partial W} = \frac{\partial f(S)}{\partial W} = \frac{\partial f(S)}{\partial S} \bullet \frac{\partial S}{\partial W} = \frac{\partial f(S)}{\partial S} \bullet \frac{\partial (Y \bullet W)}{\partial W}$$

- $f(S)$ is sigmoid, so $\frac{\partial f(S)}{\partial S} = f(S)(1 - f(S)) = z_1(1 - z_1)$, a scalar:

$$z_1(1 - z_1) \frac{\partial (Y \bullet W)}{\partial W}$$

- Using the product rule:

$$z_1(1 - z_1) \frac{\partial (Y \bullet W)}{\partial W} = z_1(1 - z_1) \left[\frac{\partial Y}{\partial W} \bullet W + Y \bullet \frac{\partial W}{\partial W} \right]$$

- Y is independent of W , so red term = 0. But $\frac{\partial W}{\partial W}$ is not simply 1, but:

$$\begin{pmatrix} \left| \begin{array}{c} 1 \\ 0 \end{array} \right| \\ \left| \begin{array}{c} 0 \\ 1 \end{array} \right| \end{pmatrix}$$

$$Y \bullet \frac{\partial W}{\partial W} = (y_1, y_2) \bullet \begin{pmatrix} \begin{vmatrix} 1 \\ 0 \end{vmatrix} \\ \begin{vmatrix} 0 \\ 1 \end{vmatrix} \end{pmatrix} = \begin{pmatrix} y_1 \\ y_2 \end{pmatrix}$$

- Putting it all together:

$$\frac{\partial Z}{\partial W} = z_1(1 - z_1) \times Y \bullet \frac{\partial W}{\partial W} = \begin{pmatrix} z_1(1 - z_1)y_1 \\ z_1(1 - z_1)y_2 \end{pmatrix} = \begin{pmatrix} \frac{\partial Z}{\partial w_{1,1}} \\ \frac{\partial Z}{\partial w_{2,1}} \end{pmatrix}$$

- For each case, c_k , in a minibatch, the forward pass will produce values for X , Y , and Z . Those values will be used to compute actual numeric gradients by filling in these symbolic gradients:

$$\frac{\partial Z}{\partial W}|_{c_k} = \begin{pmatrix} z_1(1 - z_1)y_1 \\ z_1(1 - z_1)y_2 \end{pmatrix}|_{c_k}$$

Tensor Calculations across Several Layers

Goal: Compute $\frac{\partial Z}{\partial V}$

- Expanding using the Chain Rule (with $\frac{\partial f(S)}{\partial S} = z_1(1 - z_1)$, a scalar):

$$\frac{\partial Z}{\partial V} = \frac{\partial f(S)}{\partial V} = \frac{\partial f(S)}{\partial S} \bullet \frac{\partial S}{\partial V} = z_1(1 - z_1) \frac{\partial (Y \bullet W)}{\partial V}$$

- Using the Product Rule:

$$\frac{\partial (Y \bullet W)}{\partial V} = \frac{\partial Y}{\partial V} \bullet W + Y \bullet \frac{\partial W}{\partial V} = \frac{\partial (g(R))}{\partial V} \bullet W$$

- W is indep. of V , so red term = 0; and $Y = g(R)$. Using Chain Rule:

$$\frac{\partial (g(R))}{\partial V} \bullet W = \left[\frac{\partial g(R)}{\partial R} \bullet \frac{\partial R}{\partial V} \right] \bullet W$$

- $g(R)$ is sigmoid, and R is a vector (size 2), so:

$$\frac{\partial g(R)}{\partial R} = \begin{bmatrix} \frac{\partial g(r_1)}{\partial r_1} & \frac{\partial g(r_1)}{\partial r_2} \\ \frac{\partial g(r_2)}{\partial r_1} & \frac{\partial g(r_2)}{\partial r_2} \end{bmatrix} = \begin{bmatrix} y_1(1 - y_1) & 0 \\ 0 & y_2(1 - y_2) \end{bmatrix}$$

- Since $R = X \bullet V$, and using the Product Rule:

$$\frac{\partial R}{\partial V} = \frac{\partial (X \bullet V)}{\partial V} = X \bullet \frac{\partial V}{\partial V} + \frac{\partial X}{\partial V} \bullet V$$

- X is indep. of V , so red term = 0; $\frac{\partial V}{\partial V}$ is 4-dimensional:

$$\frac{\partial V}{\partial V} = \begin{pmatrix} \begin{vmatrix} 1 & 0 \\ 0 & 0 \\ 0 & 0 \end{vmatrix} & \begin{vmatrix} 0 & 1 \\ 0 & 0 \\ 0 & 0 \end{vmatrix} \\ \begin{vmatrix} 0 & 0 \\ 1 & 0 \\ 0 & 0 \end{vmatrix} & \begin{vmatrix} 0 & 0 \\ 0 & 1 \\ 0 & 0 \end{vmatrix} \\ \begin{vmatrix} 0 & 0 \\ 0 & 0 \\ 1 & 0 \end{vmatrix} & \begin{vmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 1 \end{vmatrix} \end{pmatrix}$$

- Take dot product of $X = [x_1, x_2, x_3]$ with each internal maxtrix of $\frac{\partial V}{\partial V}$:

$$X \bullet \frac{\partial V}{\partial V} = [x_1, x_2, x_3] \bullet \begin{pmatrix} \left| \begin{array}{cc|cc} 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{array} \right| \\ \left| \begin{array}{cc|cc} 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{array} \right| \\ \left| \begin{array}{cc|cc} 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 \end{array} \right| \end{pmatrix}$$

$$= \begin{pmatrix} \left| \begin{array}{cc|cc} x_1 & 0 & 0 & x_1 \\ x_2 & 0 & 0 & x_2 \\ x_3 & 0 & 0 & x_3 \end{array} \right| \end{pmatrix}$$

Computing $\frac{\partial Z}{\partial V}$

Putting the pieces back together for $\frac{\partial Y}{\partial V} = \frac{\partial g(R)}{\partial R} \bullet \frac{\partial R}{\partial V}$:

$$\frac{\partial g(R)}{\partial R} \bullet \frac{\partial R}{\partial V} = \begin{bmatrix} y_1(1-y_1) & 0 \\ 0 & y_2(1-y_2) \end{bmatrix} \bullet \begin{pmatrix} | & x_1 & 0 & | & | & 0 & x_1 & | \\ | & x_2 & 0 & | & | & 0 & x_2 & | \\ | & x_3 & 0 & | & | & 0 & x_3 & | \end{pmatrix}$$
$$\frac{\partial Y}{\partial V} = \begin{pmatrix} | & x_1 y_1 (1-y_1) & 0 & | & | & 0 & x_1 y_2 (1-y_2) & | \\ | & x_2 y_1 (1-y_1) & 0 & | & | & 0 & x_2 y_2 (1-y_2) & | \\ | & x_3 y_1 (1-y_1) & 0 & | & | & 0 & x_3 y_2 (1-y_2) & | \end{pmatrix}$$

Note: Row vectors of $\frac{\partial R}{\partial V}$ need to be treated as column vectors for the internal matrix multiplication. The results of each multiplication are then transposed (to get row vectors). **Twist/Hack.**

$$\begin{aligned} \frac{\partial(g(R))}{\partial V} \bullet W &= \left[\frac{\partial g(R)}{\partial R} \bullet \frac{\partial R}{\partial V} \right] \bullet W \\ &= \begin{pmatrix} \begin{vmatrix} x_1 y_1 (1 - y_1) & 0 \\ x_2 y_1 (1 - y_1) & 0 \\ x_3 y_1 (1 - y_1) & 0 \end{vmatrix} & \begin{vmatrix} 0 & x_1 y_2 (1 - y_2) \\ 0 & x_2 y_2 (1 - y_2) \\ 0 & x_3 y_2 (1 - y_2) \end{vmatrix} \end{pmatrix} \bullet \begin{pmatrix} w_{11} \\ w_{21} \end{pmatrix} \\ &= \begin{pmatrix} x_1 y_1 (1 - y_1) w_{11} & x_1 y_2 (1 - y_2) w_{21} \\ x_2 y_1 (1 - y_1) w_{11} & x_2 y_2 (1 - y_2) w_{21} \\ x_3 y_1 (1 - y_1) w_{11} & x_3 y_2 (1 - y_2) w_{21} \end{pmatrix} = \frac{\partial(Y \bullet W)}{\partial V} \end{aligned}$$

Each of the 6 inner vectors is dotted with W . This is compatible with `numpy.dot` but may be a theoretical **twist/hack**.

...and finally...

$$\begin{aligned}\frac{\partial Z}{\partial V} &= \frac{\partial f(S)}{\partial V} = z_1(1 - z_1) \frac{\partial(Y \bullet W)}{\partial V} \\&= z_1(1 - z_1) \times \begin{pmatrix} x_1 y_1 (1 - y_1) w_{11} & x_1 y_2 (1 - y_2) w_{21} \\ x_2 y_1 (1 - y_1) w_{11} & x_2 y_2 (1 - y_2) w_{21} \\ x_3 y_1 (1 - y_1) w_{11} & x_3 y_2 (1 - y_2) w_{21} \end{pmatrix} \\&= \begin{pmatrix} z_1(1 - z_1) x_1 y_1 (1 - y_1) w_{11} & z_1(1 - z_1) x_1 y_2 (1 - y_2) w_{21} \\ z_1(1 - z_1) x_2 y_1 (1 - y_1) w_{11} & z_1(1 - z_1) x_2 y_2 (1 - y_2) w_{21} \\ z_1(1 - z_1) x_3 y_1 (1 - y_1) w_{11} & z_1(1 - z_1) x_3 y_2 (1 - y_2) w_{21} \end{pmatrix} \\&= \begin{pmatrix} \frac{\partial Z}{\partial v_{11}} & \frac{\partial Z}{\partial v_{12}} \\ \frac{\partial Z}{\partial v_{21}} & \frac{\partial Z}{\partial v_{22}} \\ \frac{\partial Z}{\partial v_{31}} & \frac{\partial Z}{\partial v_{32}} \end{pmatrix}\end{aligned}$$

Gradients and the Minibatch

- For every case c_k in a minibatch, the values of X , Y and Z are computed during the forward pass. The symbolic gradients in $\frac{\partial Z}{\partial V}$ are filled in using these values in X, Y and Z (along with W), producing one numeric, 3×2 matrix **per case**:

$$\left(\frac{\partial Z}{\partial V}\right)_{|c_k} = \begin{pmatrix} \frac{\partial Z}{\partial v_{11}} & \frac{\partial Z}{\partial v_{12}} \\ \frac{\partial Z}{\partial v_{21}} & \frac{\partial Z}{\partial v_{22}} \\ \frac{\partial Z}{\partial v_{31}} & \frac{\partial Z}{\partial v_{32}} \end{pmatrix}_{|c_k}$$
$$= \begin{pmatrix} z_1(1-z_1)x_1y_1(1-y_1)w_{11} & z_1(1-z_1)x_1y_2(1-y_2)w_{21} \\ z_1(1-z_1)x_2y_1(1-y_1)w_{11} & z_1(1-z_1)x_2y_2(1-y_2)w_{21} \\ z_1(1-z_1)x_3y_1(1-y_1)w_{11} & z_1(1-z_1)x_3y_2(1-y_2)w_{21} \end{pmatrix}_{|c_k}$$

Gradients and the Minibatch

- In most Deep Learning situations, the gradients will be based on a loss function, L , not simply the output of the final layer. But that's just one more level of derivative calculations (see below).
- At the completion of a minibatch, the numeric gradient matrices are added together to yield the complete gradient, which is then used to update the weights.
- For any weight matrix U in the network, and minibatch M , update weight $u_{i,j}$ as follows:


$$\left(\frac{\partial L}{\partial u_{i,j}}\right)_{|M} = \sum_{c_k \in M} \left(\frac{\partial L}{\partial u_{i,j}}\right)_{|c_k}$$

η = learning rate

$$\Delta u_{i,j} = -\eta \left(\frac{\partial L}{\partial u_{i,j}}\right)_{|M}$$

- Tensorflow and PyTorch do all of this for you.
- Everytime you write Deep Learning code, be grateful!!

The Jacobian Product Chain



$$\frac{\partial L}{\partial W} = \frac{\partial L}{\partial N} \cdot \frac{\partial N}{\partial N-1} \cdots \frac{\partial M+1}{\partial M} \cdot \frac{\partial M}{\partial W}$$

Jacobian Matrix Linking Neighboring Layers: $Y \rightarrow Z$

$$\mathbf{J}_Y^Z = \begin{pmatrix} \frac{\partial z_1}{\partial y_1} & \frac{\partial z_1}{\partial y_2} & \dots & \frac{\partial z_1}{\partial y_n} \\ \frac{\partial z_2}{\partial y_1} & \frac{\partial z_2}{\partial y_2} & \dots & \frac{\partial z_2}{\partial y_n} \\ \vdots & \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \vdots \\ \frac{\partial z_m}{\partial y_1} & \frac{\partial z_m}{\partial y_2} & \dots & \frac{\partial z_m}{\partial y_n} \end{pmatrix}$$


Note: This Jacobian is in **numerator** format.

Jacobian Linking Outputs(Z) to Incoming Weights (W)

$$\mathbf{J}_W^Z = \begin{pmatrix} \frac{\partial Z}{\partial w_{11}} & \frac{\partial Z}{\partial w_{12}} & \dots & \frac{\partial Z}{\partial w_{1m}} \\ \frac{\partial Z}{\partial w_{21}} & \frac{\partial Z}{\partial w_{22}} & \dots & \frac{\partial Z}{\partial w_{2m}} \\ \vdots & \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \vdots \\ \frac{\partial Z}{\partial w_{n1}} & \frac{\partial Z}{\partial w_{n2}} & \dots & \frac{\partial Z}{\partial w_{nm}} \end{pmatrix}$$

Note: This Jacobian is in **denominator** format (since the high-level form, i.e. the first 2 dimensions, are that of the denominator tensor), which is most convenient (readable) for weight-to-layer derivatives.

Distant Gradients = Sums of Path Products



$$\frac{\partial z_3}{\partial x_2} = \frac{\partial z_3}{\partial y_1} \frac{\partial y_1}{\partial x_2} + \frac{\partial z_3}{\partial y_2} \frac{\partial y_2}{\partial x_2} + \frac{\partial z_3}{\partial y_3} \frac{\partial y_3}{\partial x_2}$$

More generally:

$$\frac{\partial z_i}{\partial x_j} = \sum_{k \in Y} \frac{\partial z_i}{\partial y_k} \frac{\partial y_k}{\partial x_j}$$

This results from the multiplication of two Jacobian matrices.

Multiplying Jacobian Matrices

$$J_y^z \bullet J_x^y = \begin{vmatrix} \vdots & \vdots & \vdots \\ \frac{\partial z_i}{\partial y_1} & \frac{\partial z_i}{\partial y_2} & \frac{\partial z_i}{\partial y_3} \\ \vdots & \vdots & \vdots \end{vmatrix} \bullet \begin{vmatrix} \dots & \frac{\partial y_1}{\partial x_j} & \dots \\ \dots & \frac{\partial y_2}{\partial x_j} & \dots \\ \dots & \frac{\partial y_3}{\partial x_j} & \dots \end{vmatrix}$$

$$= \begin{vmatrix} \vdots & \vdots & \vdots & \vdots & \vdots \\ \dots & \dots & \frac{\partial z_i}{\partial y_1} \frac{\partial y_1}{\partial x_j} + \frac{\partial z_i}{\partial y_2} \frac{\partial y_2}{\partial x_j} + \frac{\partial z_i}{\partial y_3} \frac{\partial y_3}{\partial x_j} & \dots & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{vmatrix} = J_x^z$$

We can do this repeatedly to form J_m^n , the Jacobian relating the activation levels of upstream layer m with those of downstream layer n:

- R = Identity Matrix
- For q = n down to m+1 do:
 - $R \leftarrow R \bullet J_{q-1}^q$
- $J_m^n \leftarrow R$

Last of the Jacobians

Once we've computed J_m^n , we need one final Jacobian: J_w^m , where w are the weights feeding into layer m . Then we have the standard Jacobian, J_w^n needed for updating all weights in matrix w .

$$J_w^n \leftarrow J_m^n \bullet J_w^m$$

Weights V from layer X to Y (from the earlier example)

$$J_V^Y = \frac{\partial Y}{\partial V} = \left| \begin{array}{cc} \left(\begin{array}{cc} \frac{\partial y_1}{\partial v_{11}} & \frac{\partial y_2}{\partial v_{11}} \end{array} \right) & \left(\begin{array}{cc} \frac{\partial y_1}{\partial v_{12}} & \frac{\partial y_2}{\partial v_{12}} \end{array} \right) \\ \left(\begin{array}{cc} \frac{\partial y_1}{\partial v_{21}} & \frac{\partial y_2}{\partial v_{21}} \end{array} \right) & \dots \\ \left(\begin{array}{cc} \frac{\partial y_1}{\partial v_{31}} & \frac{\partial y_2}{\partial v_{31}} \end{array} \right) & \dots \end{array} \right|$$

- Since a) the act func for Y is sigmoid, and b) $\frac{\partial y_k}{\partial v_{ij}} = 0$ when $j \neq k$:

$$J_V^Y = \frac{\partial Y}{\partial V} = \left| \begin{array}{cc} \left(\begin{array}{cc} y_1(1-y_1)x_1 & 0 \end{array} \right) & \left(\begin{array}{cc} 0 & y_2(1-y_2)x_1 \end{array} \right) \\ \left(\begin{array}{cc} y_1(1-y_1)x_2 & 0 \end{array} \right) & \left(\begin{array}{cc} 0 & y_2(1-y_2)x_2 \end{array} \right) \\ \left(\begin{array}{cc} y_1(1-y_1)x_3 & 0 \end{array} \right) & \left(\begin{array}{cc} 0 & y_2(1-y_2)x_3 \end{array} \right) \end{array} \right|$$

Last of the Jacobian Multiplications

From the previous example (the network with layer sizes 3-2-1), once we have J_Y^Z and J_V^Y , we can multiply (taking dot products of J_Y^Z with the inner vectors of J_V^Y) to produce J_V^Z : the matrix of gradients that backprop uses to modify the weights in V.



$$J_V^Z = J_Y^Z \bullet J_V^Y = \begin{vmatrix} \frac{\partial z_1}{\partial y_1} & \frac{\partial z_1}{\partial y_2} \end{vmatrix} \bullet \frac{\partial Y}{\partial V}$$

- Since a) the act func for Z is sigmoid, and b) $\frac{\partial \text{sum}(z_k)}{\partial y_i} = w_{ik}$ (where $\text{sum}(z_k)$ = sum of weighted inputs to node z_k):

$$= \begin{vmatrix} z_1(1-z_1)w_{11} & z_1(1-z_1)w_{21} \end{vmatrix} \bullet \frac{\partial Y}{\partial V}$$

$$J_V^Z = \begin{pmatrix} z_1(1-z_1)x_1y_1(1-y_1)w_{11} & z_1(1-z_1)x_1y_2(1-y_2)w_{21} \\ z_1(1-z_1)x_2y_1(1-y_1)w_{11} & z_1(1-z_1)x_2y_2(1-y_2)w_{21} \\ z_1(1-z_1)x_3y_1(1-y_1)w_{11} & z_1(1-z_1)x_3y_2(1-y_2)w_{21} \end{pmatrix}$$

One entry of J_V^Z



$$J_V^Z(1, 1) = z_1(1 - z_1)x_1y_1(1 - y_1)w_{11}$$

First of the Jacobians

- This iterative process is very general, and permits the calculation of J_M^N ($M < N$) for any layers M and N, or J_W^N for an weights (W) upstream of N.
- However, the standard situation in backpropagation is to compute $J_{W_i}^L \forall i$ where L is the objective (loss) function and W_i are the weight matrices.
- This follows the same procedure as sketched above, but the series of dot products **begins** with J_N^L : the Jacobian of derivatives of the loss function w.r.t. the activations of the output layer.
- For example, assume L = Mean Squared Error (MSE), Z is an output layer of 3 sigmoid nodes, and t_i are the target values for those 3 nodes for a particular case (c):

$$L(c) = \frac{1}{3} \sum_{i=1}^3 (z_i - t_i)^2$$

Taking partial derivatives of L(c) w.r.t. the z_i yields:

$$J_Z^L = \frac{\partial L}{\partial Z} = \left[\frac{2}{3}(z_1 - t_1) \quad \frac{2}{3}(z_2 - t_2) \quad \frac{2}{3}(z_3 - t_3) \right]$$

First of the Jacobian Multiplications

- Continuing the example: Assume that layer Y (of size 2) feeds into layer Z, using weights W, then:

$$J_Y^Z = \begin{vmatrix} z_1(1-z_1)w_{11} & z_1(1-z_1)w_{21} \\ z_2(1-z_2)w_{12} & z_2(1-z_2)w_{22} \\ z_3(1-z_3)w_{13} & z_3(1-z_3)w_{23} \end{vmatrix}$$

- Our first Jacobian multiplication yields J_Y^L , a 1 x 2 row vector (shown transposed to fit on the page):

$$J_Z^L \bullet J_Y^Z = J_Y^L = \begin{vmatrix} \frac{\partial L}{\partial y_1} & \frac{\partial L}{\partial y_2} \end{vmatrix} =$$

$$\begin{vmatrix} \frac{2}{3}(z_1 - t_1)z_1(1-z_1)w_{11} + \frac{2}{3}(z_2 - t_2)z_2(1-z_2)w_{12} + \frac{2}{3}(z_3 - t_3)z_3(1-z_3)w_{13} \\ \frac{2}{3}(z_1 - t_1)z_1(1-z_1)w_{21} + \frac{2}{3}(z_2 - t_2)z_2(1-z_2)w_{22} + \frac{2}{3}(z_3 - t_3)z_3(1-z_3)w_{23} \end{vmatrix}^T$$


Backpropagation with Tensors: The Big Picture

General Algorithm

- Assume Layer M is upstream of Layer N (the output layer). So $M < N$.
- Assume V is the tensor of weights feeding into Layer M.
- Assume L is the loss function.
- Goal: Compute $J_V^L = \frac{\partial L}{\partial V}$
- $R = J_N^L$ (the partial derivatives of the loss function w.r.t. the output layer)
- If output layer is Softmax: $R \leftarrow R \bullet J^{soft}$
- For $q = N$ down to $M + 1$ do:
 - $R \leftarrow R \bullet J_{q-1}^q$
- $J_V^L \leftarrow R \bullet J_V^M$
- Use J_V^L to update the weights in V.

... And now some practical details on implementing this...

Building the J_Y^Z and J_W^Z Jacobians



Detailed Entries of J_Y^Z

$$J_Y^Z =$$

$$\begin{pmatrix} w_{11}z_1(1-z_1) & w_{21}z_1(1-z_1) & \cdots & w_{n1}z_1(1-z_1) \\ w_{12}z_2(1-z_2) & w_{22}z_2(1-z_2) & \cdots & w_{n2}z_2(1-z_2) \\ \vdots & \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \vdots \\ w_{1m}z_m(1-z_m) & w_{2m}z_m(1-z_m) & \cdots & w_{nm}z_m(1-z_m) \end{pmatrix}$$

More succinctly:

$$J_Y^Z = (W \bullet J_{Sum}^Z)^T = J_{Sum}^Z \bullet W^T$$

Jacobian J_{Sum}^Z Linking Z to Summed Inputs

$$J_{Sum}^Z =$$

$$\begin{pmatrix} z_1(1 - z_1) & 0 & \cdots & 0 \\ 0 & z_2(1 - z_2) & \cdots & 0 \\ \vdots & \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \vdots \\ 0 & 0 & \cdots & z_m(1 - z_m) \end{pmatrix}$$

Simplifying J_W^Z

In an earlier slide, J_W^Z was presented. Note that it has the same shape as the weight matrix, W :

$$J_W^Z = \begin{pmatrix} \frac{\partial Z}{\partial w_{11}} & \frac{\partial Z}{\partial w_{12}} & \cdots & \frac{\partial Z}{\partial w_{1m}} \\ \frac{\partial Z}{\partial w_{21}} & \frac{\partial Z}{\partial w_{22}} & \cdots & \frac{\partial Z}{\partial w_{2m}} \\ \vdots & \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \vdots \\ \frac{\partial Z}{\partial w_{n1}} & \frac{\partial Z}{\partial w_{n2}} & \cdots & \frac{\partial Z}{\partial w_{nm}} \end{pmatrix}$$

- Weight w_{ik} connects y_i to z_k .
- So z_k is the only element of Z that w_{ik} affects.
- Thus, $\frac{\partial Z}{\partial w_{ik}} = [0, \dots, y_i z_k (1 - z_k), \dots, 0]^T$ (only k th entry is non-zero).
- Make a simpler, more practical, matrix of only these positive values: \hat{J}_W^Z .

The Simplified Jacobian: \hat{J}_W^Z


$$\begin{pmatrix} \frac{\partial z_1}{\partial w_{11}} & \frac{\partial z_2}{\partial w_{12}} & \cdots & \frac{\partial z_m}{\partial w_{1m}} \\ \frac{\partial z_1}{\partial w_{21}} & \frac{\partial z_2}{\partial w_{22}} & \cdots & \frac{\partial z_m}{\partial w_{2m}} \\ \vdots & \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \vdots \\ \frac{\partial z_1}{\partial w_{n1}} & \frac{\partial z_2}{\partial w_{n2}} & \cdots & \frac{\partial z_m}{\partial w_{nm}} \end{pmatrix} = \begin{pmatrix} y_1 z_1 (1 - z_1) & y_1 z_2 (1 - z_2) & \cdots & y_1 z_m (1 - z_m) \\ y_2 z_1 (1 - z_1) & y_2 z_2 (1 - z_2) & \cdots & y_2 z_m (1 - z_m) \\ \vdots & \vdots & \vdots & \vdots \\ \vdots & \vdots & \vdots & \vdots \\ y_n z_1 (1 - z_1) & y_n z_2 (1 - z_2) & \cdots & y_n z_m (1 - z_m) \end{pmatrix}$$

More succinctly:

$$\hat{J}_W^Z = Y \otimes \text{Diag}(J_{\text{Sum}}^Z)$$

where \otimes = Outer Product and $\text{Diag}(J_{\text{Sum}}^Z)$ = diagonal of J_{Sum}^Z .

The Backward Pass across a Layer



- 1 Receive J_Z^L from downstream.
- 2 Compute $J_W^L \leftarrow J_Z^L \bullet J_W^Z$ and use it to update W .
- 3 Compute $J_Y^L \leftarrow J_Z^L \bullet J_Y^Z$ and pass it back (upstream).

Details of $J_W^L \leftarrow J_Z^L \bullet J_W^Z$


- Remember that every element of J_W^Z is a vector, so $J_Z^L \bullet J_W^Z$ involves many vector dot products: $J_Z^L \bullet Q$ for each vector (Q) in J_W^Z . This produces a final matrix with the same dimensions as W.
- In numpy, the "dot" function produces the desired result, but it is sensitive to argument order; the proper call is: `numpy.dot(J_W^Z , J_Z^L)`
- If we use \hat{J}_W^Z instead of J_W^Z , then the dot product is replaced by this:

$$J_W^L \leftarrow J_Z^L \times \hat{J}_W^Z$$

- The use of "×" means that the ith element of J_Z^L is multiplied by every item in the ith column of \hat{J}_W^Z
- In numpy, the standard multiplication operator, "*" between the vector J_Z^L and the matrix \hat{J}_W^Z will perform the desired operation, as long as $\|J_Z^L\| =$ the number of columns in \hat{J}_W^Z .

SoftMax: One More Layer (but without weights)

$$\text{Softmax}(Z) = S(Z)$$



$$s_i = S(z_i) = \frac{e^{z_i}}{\sum_{z_k \in Z} e^{z_k}} = \frac{e^{z_i}}{\Sigma}$$

Derivatives of Softmax

Effect of z_i on $S(z_i)$

$$\begin{aligned}\frac{\partial S(z_i)}{\partial z_i} &= \frac{\frac{\partial e^{z_i}}{\partial z_i} \Sigma - \frac{\partial \Sigma}{\partial z_i} e^{z_i}}{\Sigma^2} = \frac{e^{z_i} \Sigma - e^{z_i} e^{z_i}}{\Sigma^2} = \\ \frac{e^{z_i}}{\Sigma} - \left(\frac{e^{z_i}}{\Sigma} \right)^2 &= S(z_i) - S(z_i)^2 = s_i - s_i^2 = s_i(1 - s_i)\end{aligned}$$

Effect of z_i on $S(z_k)$ where $i \neq k$


$$\begin{aligned}\frac{\partial S(z_k)}{\partial z_i} &= \frac{\frac{\partial e^{z_k}}{\partial z_i} \Sigma - \frac{\partial \Sigma}{\partial z_i} e^{z_k}}{\Sigma^2} = \frac{0 - e^{z_i} e^{z_k}}{\Sigma^2} = \\ -\frac{e^{z_i}}{\Sigma} \frac{e^{z_k}}{\Sigma} &= -S(z_i)S(z_k) = -s_i s_k\end{aligned}$$

The Softmax ($m \times m$) Jacobian Matrix

J^{soft}

$$\mathbf{J}_Z^S = \begin{pmatrix} \frac{\partial s_1}{\partial z_1} & \frac{\partial s_1}{\partial z_2} & \cdots & \frac{\partial s_1}{\partial z_m} \\ \frac{\partial s_2}{\partial z_1} & \frac{\partial s_2}{\partial z_2} & \cdots & \frac{\partial s_2}{\partial z_m} \\ \vdots & \vdots & \vdots & \vdots \\ \frac{\partial s_m}{\partial z_1} & \frac{\partial s_m}{\partial z_2} & \cdots & \frac{\partial s_m}{\partial z_m} \end{pmatrix} = \begin{pmatrix} s_1 - s_1^2 & -s_1 s_2 & \cdots & -s_1 s_m \\ -s_2 s_1 & s_2 - s_2^2 & \cdots & -s_2 s_m \\ \vdots & \vdots & \vdots & \vdots \\ -s_m s_1 & -s_m s_2 & \cdots & s_m - s_m^2 \end{pmatrix}$$

Backward Pass Across Softmax



$$J_Z^L = J_S^L \bullet J_Z^S$$


This assumes J_S^L is a row vector. If it's a column vector, then:

$$J_Z^L = (J_S^L)^T \bullet J_Z^S$$

Practical Tips

- 1 Only add as many hidden layers and hidden nodes as necessary. Too many \rightarrow more weights to learn + increased chance of over-specialization.
- 2 Scale all input values to the same range, typically $[0\ 1]$ or $[-1\ 1]$.
- 3 Use target values of 0.1 (for zero) and 0.9 (for 1) to avoid saturation effects of sigmoids.
- 4 Beware of tricky encodings of input (and decodings of output) values. Don't combine too much info into a single node's activation value (even though it's fun to try), since this can make proper weights difficult (or impossible) to learn.
- 5 For discrete (e.g. nominal) values, one (input or output) node per value is often most effective. E.g. car model and city of residence -vs- income and education for assessing car-insurance risk.
- 6 All initial weights should be relatively small: $[-0.1\ 0.1]$
- 7 Bias nodes can be helpful for complicated data sets.
- 8 Check that all your layer sizes, activation functions, activation ranges, weight ranges, learning rates, etc. *make sense* in terms of each other and your goals for the ANN. One improper choice can ruin the results.

Supervised Learning in the Cerebellum



- Granular cells detect contexts.
- Parallel fibers and Purkinje cells map contexts to actions
- Climbing fibers from Inferior Olive provide (supervisory) feedback signals for LTD on Parallel-Purkinje synapses